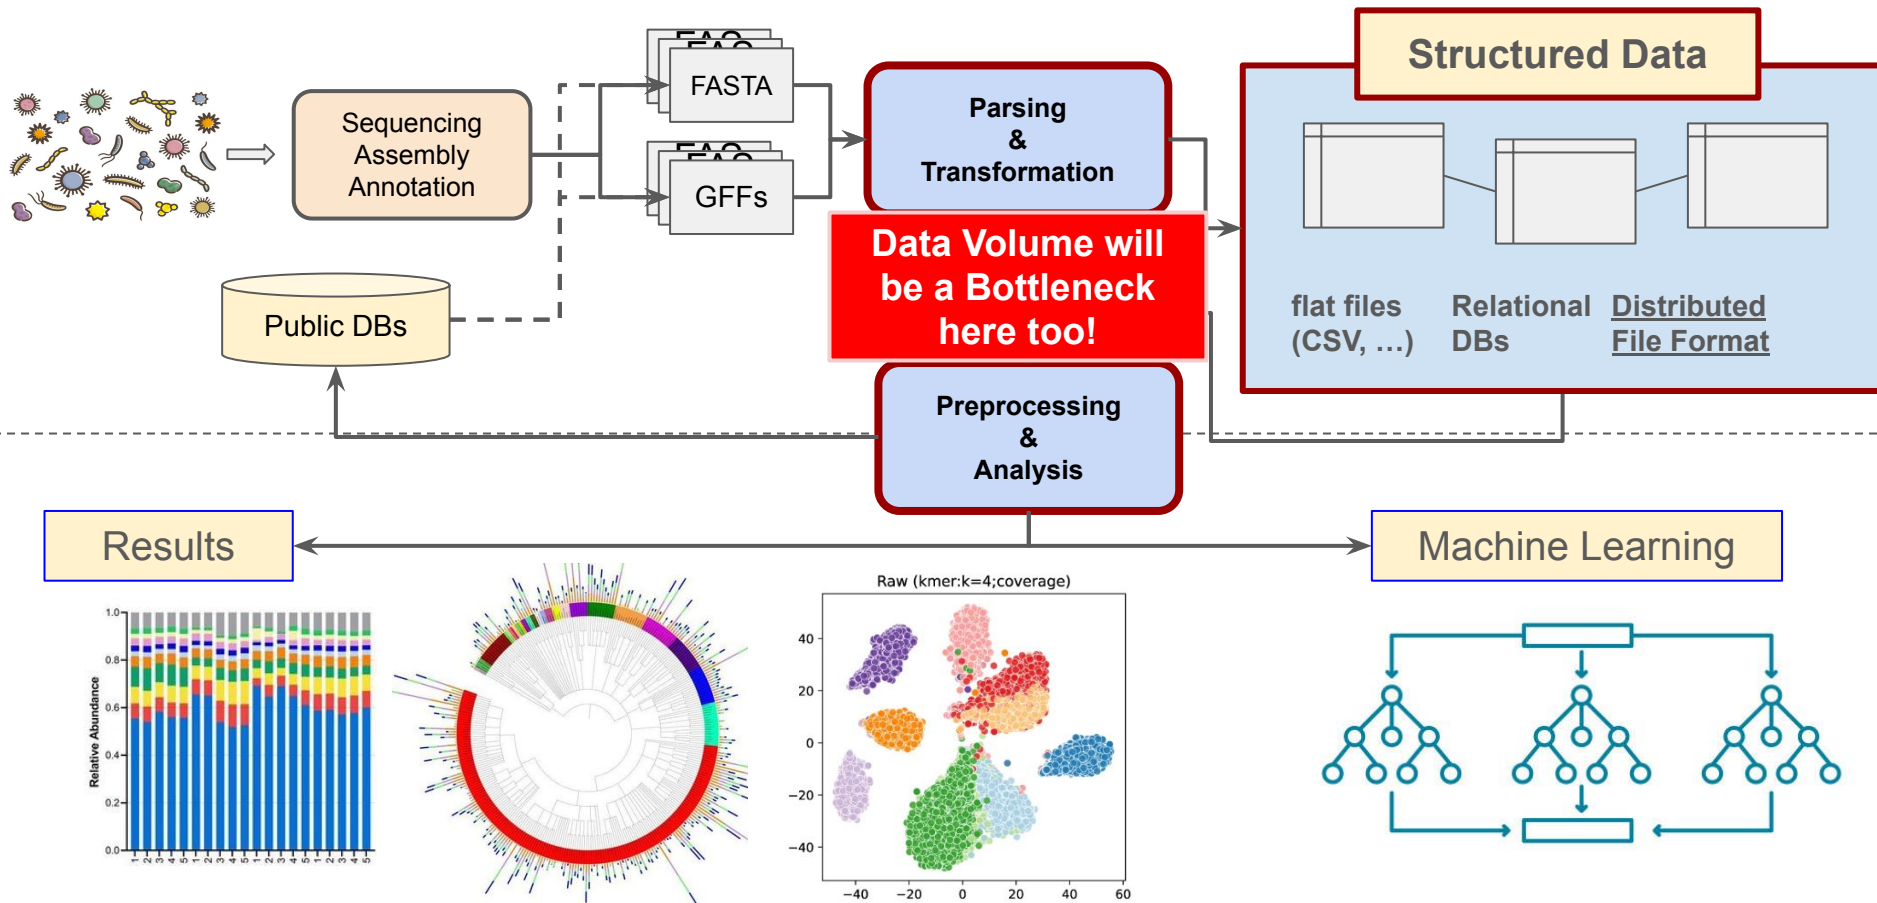


Scalable [Meta]genomic Data Analysis on Apache Spark

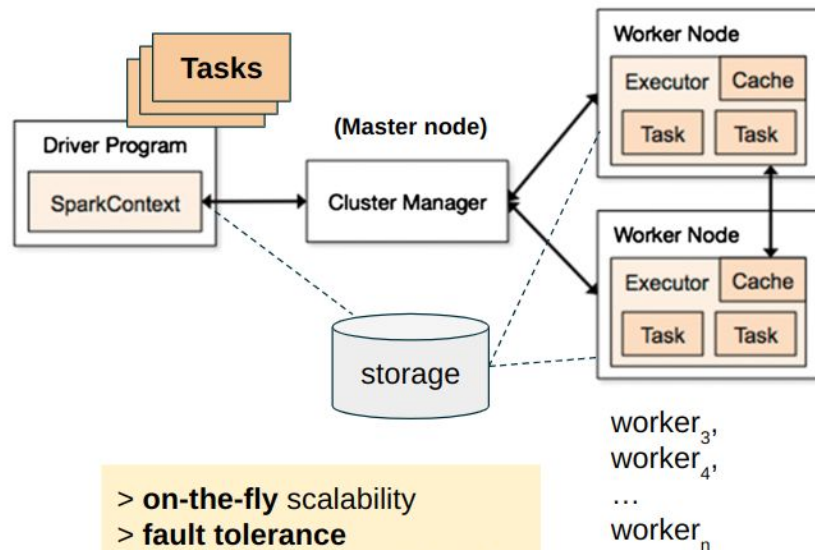
Satria Kautsar
Genome Analysis R&D scientist
October 4, 2024



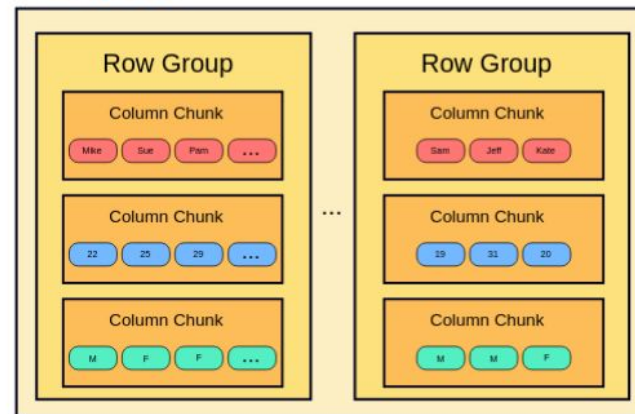
What's after assembly & annotation?



The solution: process & data scalability



- > **on-the-fly** scalability
- > **fault tolerance**
- > automated task optimization




- > data is **partitioned** by **row groups**, then sorted by **column chunks**
- > each chunk is **indexed** (including min, max values)
- > resulted in a **distributable, compressible** data

The thing is...

Working with large-scale genomics data can be **messy**

- many **data formats**
- non-computer-friendly data formats
- non-compliant data entry
- data is increasing in an **exponential** fashion

Writing and troubleshooting a genome analysis pipeline can become **a nightmare**



<https://github.com/zhongwang/axolotl>

Genome Analysis **library** written in **Python** (on top of the **pySpark** library)

Defines a collection of **Standardized Tabular Data Structures** (i.e., Dataframes with schema) for all sorts of genomic data types (not file formats): NuclSeqDF, ProtSeqDF, cdsDF, etc...

Defines a collection of **Parser Classes** to process raw genomic files (FASTA, GFF, etc) into data tables (stored as Parquet)

Defines commonly used **functions** and **workflows** to preprocess, query, and analyze genomic data

Demonstrating Axolotl



Dataset:

1,222,123 genomic (FASTA + GFF) files (bacteria, fungi, archaea) [8/22/2023]

- Originally sourced from NCBI (Genbank+Refseq) and IMG
- GFFs include CDS, **BGC** (antiSMASH + EMERALD) and other annotations
- Total size: FASTA = **5,017 GiB**; GFFs = **8,285 GiB**

>11,000,000 BGCs

(**10 times** the size of
state-of-the-art global analysis
[Gavriilidou, 2022])

Tasks:


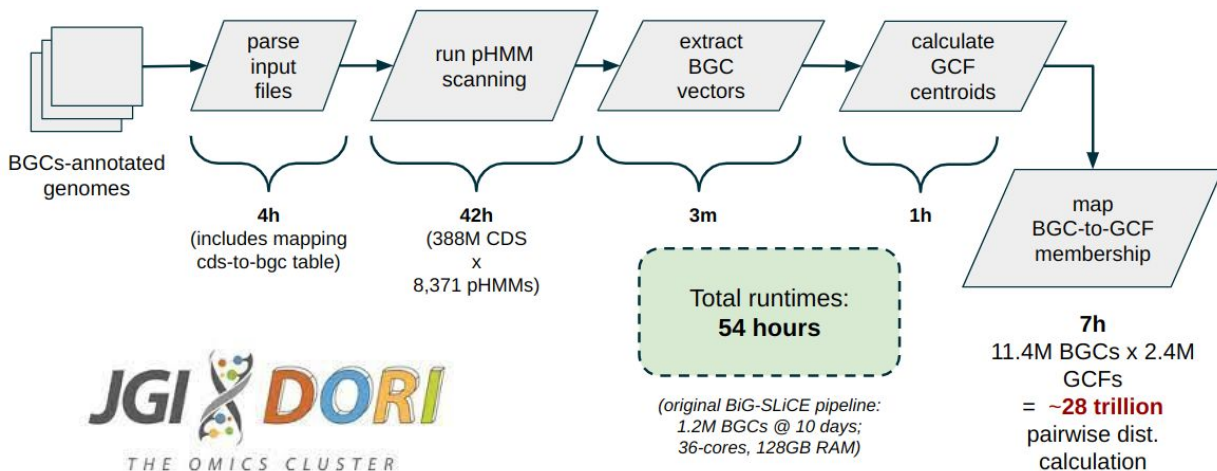
- **Parse** the raw input files and save as **Axolotl Tables** (stored in Parquet)
 - contigs
 - features => CDS, BGC
- Perform **clustering analysis** on the BGCs by replicating the "BiG-SLiCE" pipeline
- **Query** the Parquet "database" and generate some statistical summaries

Running the full "BiG-SLiCE" pipeline

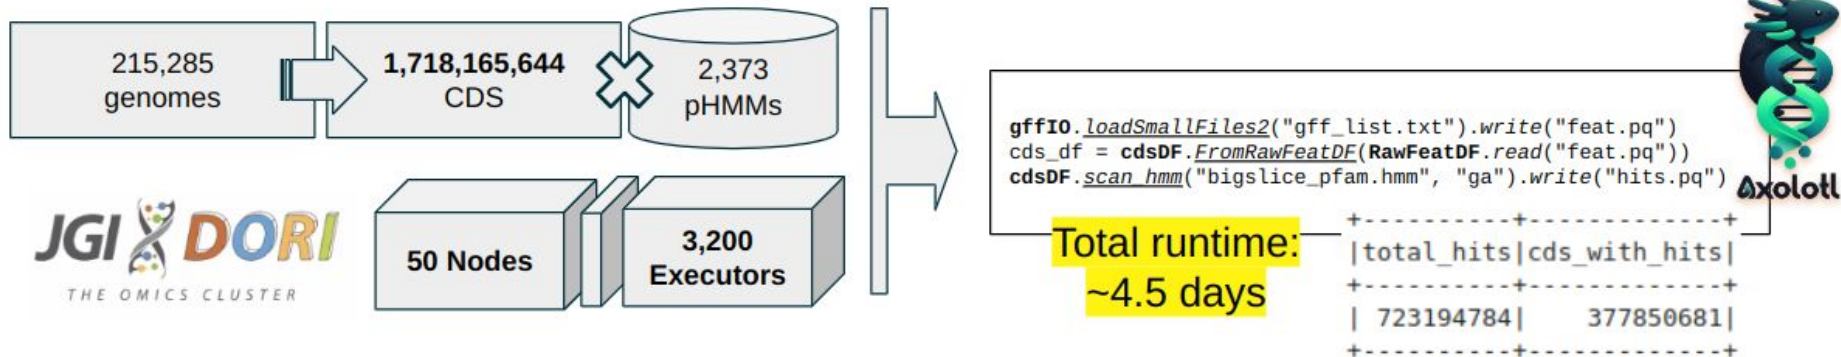
Total resources used: 16 normal Dori nodes

- 1x 16-cores, 32GB Node [Master & Driver]

- 63x 16-cores, 128GB Nodes [Workers]



We survey genome-wide 'biosynthetic' genes of all RefSeq genomes...



Once you got your tables, the real fun begins...

```
%time spark.sql("SELECT count(idx) FROM cds WHERE aa_sequence like '%FIRVEFQ%').show()
```

```
[Stage 177:=====> (1944 + 72) / 2016]
```

```
+-----+  
|count(idx)|  
+-----+  
|      14526|  
+-----+
```

CPU times: user 14.8 ms, sys: 2.66 ms, total: 17.5 ms
Wall time: 16.8 s

Sifting through 388,075,840 BGC's CDS and perform an AA substring match takes 17 seconds

```
query = (  
    "SELECT "  
        "sum(less_than_10k) as less_than_10k"  
        ",sum(10k_to_100k) as 10k_to_100k"  
        ",sum(100k_to_200k) as 100k_to_200k"  
        ",sum(more_than_200k) as more_than_200k"  
        ",count(less_than_10k) as all"  
    " FROM (SELECT "  
        "(CASE WHEN len_nt < 10000 THEN 1 ELSE 0 END) as less_than_10k"  
        ",(CASE WHEN 10000 <= len_nt AND len_nt <= 100000 THEN 1 ELSE 0 END) as 10k_to_100k"  
        ",(CASE WHEN 100000 <= len_nt AND len_nt <= 200000 THEN 1 ELSE 0 END) as 100k_to_200k"  
        ",(CASE WHEN len_nt > 200000 THEN 1 ELSE 0 END) as more_than_200k"  
    " FROM (SELECT (bgc.location.end - bgc.location.start) as len_nt FROM bgc)"  
    " ) "  
)
```

Spark supports SQL-like complex queries

*with Axolotl, working with 10B CDSes will feels like working with 10K ones!
(just put thousand CPUs on it)*

```
(381 + 614) / 995]
```

```
+-----+  
|re_than_200k|    all|  
+-----+  
|      1887|11359779|  
+-----+
```

: 2.87 ms

```
%time spark.sql(\  
    "SELECT cds.seq_id, cds.locus_tag, cds.aa_sequence" \  
    " FROM cds JOIN pfam_hits ON cds.idx=pfam_hits.cds_id" \  
    " WHERE pfam_hits.hmm_acc like 'PF01832.23'" \  
).rdd.map(\  
    lambda row: ">{}{}{}\n{}\n".format(\  
        row.seq_id, row.locus_tag, row.aa_sequence \  
    ) \  
).saveAsTextFile("./all_cds_pfam_hits-pf01832.fa")
```

CPU times: user 121 ms, s
Wall time: 2min 23s

can also be used to generate data for follow-up analyses/processing


```
>CP065473.1|ctg1_287  
MANTDFIKEIAGDAQIYKORCEASTDQUCESAGSCELSADMDGEMKGTUQVFFHQ  
YKLTQYDITTEGVPDEPSPNTPDPEPDIPSEYDGDITLNOKLKDVYFPOLHVSSODGNQVETI  
NHVYETISGTLTVRKMGLDTLKGTFPSYIFKDKSEFQSVQENFGDKFANELISEIVEDGLELD  
GMPRWAEPLRDERYKKESSMEALKKYVNPYPKMEIFDVVEYIYEPKLEITQDQFWKGTDLHLVADT  
>MHC01000076.1|B8T97_18990  
MINNSNDIGFTIDTAGLKLQKAVNGDENAGOSALTAARQFESIFTSPMLKSMRDANSDFKSLDM  
LRKTOAVOSTOFDSRHSFVTKLPYPADKAARMLGVDSSLLIAQAALETGWGQKMYKNARGNSNNLFN  
PQYADKVLRVKAIDQMLLQ  
>AAAQXU010000001.1|ctg1_50  
MINKQMKIIVMIPMLVPMYGLTTVGGQLQDSLTGENSFVKEVEAATASQQAIFDKIAAPAASQE  
VNGTSWNKDLVKKVVDATDYKVAAMELOKAGYATSPYASGLTOVENYDLAKYDVLVDKILTKST  
VNVGRAKITSPVSNIGWSKPYNVYGRFVTNATTYAQGEIKLLREAGTKGYQFSINNKITIGWID  
GWLDRNATITLYDQEEYKNTVAIDAVKVNKGNVAWTEPYRTVGTGLTGAETYLNKEVEVREAKTP  
TSRGTYEFSVDGKVIQMLDKKAFDYVDINYNKAVNLDAVEINVTGNNAVTPAYKSGKVLKTSAA  
ATAVAKTUDITDEAKTDTATVVRVUNCKTUGUARDATETRE
```


Acknowledgements



U.S. DEPARTMENT OF
ENERGY

Office of
Science



Zhong Wang
Harrison Ho
Bryce Foster
Fengchen Liu





Hands-on Workshop on Axolotl Library & Apache Spark

Satria Kautsar
Genome Analysis R&D scientist
October 4, 2024

What to expect:

You will:

- **Get introduced to Apache Spark and its primary components**
- **Walk through example how to parse, process and analyze FASTA+GFFs data using Spark and Axolotl**
- **Small exercises to keep you engaged**
- **Chance to explore and ask questions!**

- Login and create a Dataproc Cluster (covered by Steven)
- Open up the JupyterLab Interface
- Open a new Terminal screen

```
> cd /  
> git clone https://github.com/JGI-Bioinformatics/jgi-scalable-toolkit-workshop-24.git  
> gsutil cp gs://zw_axolotl/jgi_workshop_2024/data/phmms.zip /phmms.zip  
> unzip phmms.zip
```